

11/2017

Datacamp - Machine Learning with the Experts: School Budgets (Data Scientist Track with Python)

[Machine Learning with the Experts: School Budgets](#)

Course Description

Data science isn't just for predicting ad-clicks-it's also useful for social impact! This course is a case study from a machine learning competition on DrivenData. You'll explore a problem related to school district budgeting. By building a model to automatically classify items in a school's budget, it makes it easier and faster for schools to compare their spending with other schools. In this course, you'll begin by building a baseline model that is a simple, first-pass approach. In particular, you'll do some natural language processing to prepare the budgets for modeling. Next, you'll have the opportunity to try your own techniques and see how they compare to participants from the competition. Finally, you'll see how the winner was able to combine a number of expert techniques to build the most accurate model.

Part 1 : Exploring the raw data

In this chapter, you'll be introduced to the problem you'll be solving in this course. How do you accurately classify line-items in a school budget based on what that money is being used for? You will explore the raw text and numeric values in the dataset, both quantitatively and visually. And you'll learn how to measure success when trying to predict class labels for each row of the dataset.

What category of problem is this?

You're no novice to data science, but let's make sure we agree on the basics.

As Peter from DrivenData explained in the video, you're going to be working with school district budget data. This data can be classified in many ways according to certain labels, e.g. Function: Career & Academic Counseling, or Position_Type: Librarian.

Your goal is to develop a model that predicts the probability for each possible label by relying on some correctly labeled examples.

What type of machine learning problem is this?

Possible Answers : => 4

- Reinforcement Learning, because the model is learning from the data through a system of rewards and punishments.
- Unsupervised Learning, because the model doesn't output labels with certainty.
- Unsupervised Learning, because not all data is correctly classified to begin with.
- Supervised Learning, because the model will be trained using labeled examples.

Results :

Yes! Using correctly labeled budget line items to train means this is a supervised learning problem.

What is the goal of the algorithm?

As you know from previous courses, there are different types of supervised machine learning problems. In this exercise you will tell us what type of supervised machine learning problem this is, and why you think so.

Remember, your goal is to correctly label budget line items by training a supervised model to predict the probability of each possible label, taking most probable label as the correct label.

Possible Answers : => 2

- Regression, because the model will output probabilities.
- Classification, because predicted probabilities will be used to select a label class.
- egression, because probabilities take a continuous value between 0 and 1.
- Classification, because the model will output probabilities.

Results :

Nice! Specifically, we have ourselves a multi-class-multi-label classification problem (quite a mouthful!), because there are 9 broad categories that each take on many possible sub-label instances.

Loading the data

Now it's time to check out the dataset! You'll use pandas (which has been pre-imported as `pd`) to load your data into a `DataFrame` and then do some Exploratory Data Analysis (EDA) of it.

The training data is available as `TrainingData.csv`. Your first task is to load it into a `DataFrame` in the IPython Shell using `pd.read_csv()` along with the keyword argument `index_col=0`.

Use methods such as `.info()`, `.head()`, and `.tail()` to explore the budget data and the properties of the features and labels.

Some of the column names correspond to features - descriptions of the budget items - such as the `Job_Title_Description` column. The values in this column tell us if a budget item is for a teacher, custodian, or other employee.

Some columns correspond to the budget item labels you will be trying to predict with your model. For example, the `Object_Type` column describes whether the budget item is related classroom supplies, salary, travel expenses, etc.

Use `df.info()` in the IPython Shell to answer the following questions:

- How many rows are there in the training data?
- How many columns are there in the training data?
- How many non-null entries are in the `Job_Title_Description` column?

Possible Answers : => 3

- 25 rows, 1560 columns, 1560 non-null entries in `Job_Title_Description`.
- 225 rows, 1560 columns, 1131 non-null entries in `Job_Title_Description`.
- 21560 rows, 25 columns, 1131 non-null entries in `Job_Title_Description`.
- 21560 rows, 25 columns, 1560 non-null entries in `Job_Title_Description`.

Results :

```
In [4]: import pandas as pd
```

```
In [5]: df = pd.read_csv('TrainingData.csv',index_col=0)
```

```
In [6]: df.shape
```

```
Out[6]: (1560, 25)
```

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 1560 entries, 198 to 101861
```

```
Data columns (total 25 columns):
```

Function	1560 non-null object
Use	1560 non-null object
Sharing	1560 non-null object
Reporting	1560 non-null object
Student_Type	1560 non-null object
Position_Type	1560 non-null object
Object_Type	1560 non-null object
Pre_K	1560 non-null object
Operating_Status	1560 non-null object
Object_Description	1461 non-null object
Text_2	382 non-null object
SubFund_Description	1183 non-null object
Job_Title_Description	1131 non-null object
Text_3	677 non-null object
Text_4	193 non-null object
Sub_Object_Description	364 non-null object
Location_Description	874 non-null object
FTE	449 non-null float64
Function_Description	1340 non-null object
Facility_or_Department	252 non-null object
Position_Extra	1026 non-null object
Total	1542 non-null float64
Program_Description	1192 non-null object
Fund_Description	819 non-null object
Text_1	1132 non-null object

```
dtypes: float64(2), object(23)
```

```
memory usage: 316.9+ KB
```

```
In [8]: df.describe()
```

```
Out[8]:
```

	FTE	Total
count	449.000000	1.542000e+03
mean	0.493532	1.446867e+04
std	0.452844	7.916752e+04
min	-0.002369	-1.044084e+06
25%	NaN	NaN
50%	NaN	NaN
75%	NaN	NaN
max	1.047222	1.367500e+06

```
In [9]: df.head()
```

```
Out[9]:
```

	Function	Use	Sharing	Reporting \
198	NO_LABEL	NO_LABEL	NO_LABEL	NO_LABEL
209	Student Transportation	NO_LABEL	Shared Services	Non-School
750	Teacher Compensation	Instruction	School Reported	School
931	NO_LABEL	NO_LABEL	NO_LABEL	NO_LABEL
1524	NO_LABEL	NO_LABEL	NO_LABEL	NO_LABEL

	Student_Type	Position_Type	Object_Type	Pre_K \
198	NO_LABEL	NO_LABEL	NO_LABEL	NO_LABEL
209	NO_LABEL	NO_LABEL	Other Non-Compensation	NO_LABEL
750	Unspecified	Teacher	Base Salary/Compensation	Non PreK
931	NO_LABEL	NO_LABEL	NO_LABEL	NO_LABEL
1524	NO_LABEL	NO_LABEL	NO_LABEL	NO_LABEL

	Operating_Status	Object_Description \
198	Non-Operating	Supplemental *
209	PreK-12 Operating	REPAIR AND MAINTENANCE SERVICES
750	PreK-12 Operating	Personal Services - Teachers
931	Non-Operating	General Supplies
1524	Non-Operating	Supplies and Materials

	...	Sub_Object_Description \
198	...	Non-Certificated Salaries And Wages
209	...	NaN
750	...	NaN
931	...	General Supplies
1524	...	Supplies And Materials

	Location_Description	FTE	Function_Description \
198	NaN	NaN	Care and Upkeep of Building Services
209	ADMIN. SERVICES	NaN	STUDENT TRANSPORT SERVICE
750	NaN	1.0	NaN
931	NaN	NaN	Instruction
1524	NaN	NaN	Other Community Services *

	Facility_or_Department	Position_Extra	Total \
198		NaN	NaN -8291.86
209		NaN	NaN 618.29
750		NaN	TEACHER 49768.82
931	Instruction And Curriculum		NaN -1.02
1524		NaN	NaN 2304.43

	Program_Description \
198	NaN
209	PUPIL TRANSPORTATION
750	Instruction - Regular
931	"Title I, Part A Schoolwide Activities Related...
1524	NaN

	Fund_Description	Text_1
198	Title I - Disadvantaged Children/Targeted Assi...	TITLE I CARRYOVER
209	General Fund	NaN
750	General Purpose School	NaN
931	General Operating Fund	NaN
1524	Title I - Disadvantaged Children/Targeted Assi...	TITLE I PI+HOMELESS

[5 rows x 25 columns]

In [10]: df.tail()

Out[10]:

	Function	Use	Sharing \
344986	Substitute Compensation	Instruction	School Reported
384803	NO_LABEL	NO_LABEL	NO_LABEL
224382	Substitute Compensation	Instruction	School Reported
305347	Facilities & Maintenance	O&M Leadership & Management	
101861	Teacher Compensation	Instruction	School Reported

	Reporting	Student_Type	Position_Type \
344986	School	Unspecified	Substitute
384803	NO_LABEL	NO_LABEL	NO_LABEL
224382	School	Special Education	Substitute
305347	Non-School	Gifted	Custodian
101861	School	Poverty	Teacher

	Object_Type	Pre_K	Operating_Status \
344986	Benefits	NO_LABEL	PreK-12 Operating
384803	NO_LABEL	NO_LABEL	Non-Operating
224382	Substitute Compensation	NO_LABEL	PreK-12 Operating
305347	Other Compensation/Stipend	Non PreK	PreK-12 Operating
101861	Base Salary/Compensation	NO_LABEL	PreK-12 Operating

	Object_Description \
344986	EMPLOYEE BENEFITS
384803	EMPLOYEE BENEFITS
224382	OTHER PERSONAL SERVICES
305347	Extra Duty Pay/Overtime For Support Personnel
101861	SALARIES OF REGULAR EMPLOYEES

	...	\
344986	...	
384803	...	
224382	...	
305347	...	
101861	...	

	Sub_Object_Description	Location_Description \
344986		NaN NaN
384803		NaN PERSONNEL-PAID LEAVE
224382		NaN School
305347	Extra Duty Pay/Overtime For Support Personnel	Unallocated
101861		NaN NaN

FTE	Function_Description	Facility_or_Department \
344986 NaN	UNALLOC BUDGETS/SCHOOLS	NaN
384803 NaN	NON-PROJECT	NaN
224382 0.0	EXCEPTIONAL	NaN
305347 NaN	Facilities Maintenance And Operations	Gifted And Talented
101861 NaN	TITLE I	NaN

	Position_Extra	Total \
344986	PROFESSIONAL-INSTRUCTIONAL	27.04000
384803	PROFESSIONAL-INSTRUCTIONAL	NaN
224382	NaN	200.39000
305347	ANY CUS WHO IS NOT A SUPER	5.29000
101861	PROFESSIONAL-INSTRUCTIONAL	1575.03504

	Program_Description	Fund_Description \
344986	GENERAL HIGH SCHOOL EDUCATION	NaN
384803	STAFF SERVICES	NaN
224382	NaN GENERAL FUND	

305347	Gifted And Talented	General Operating Fund
101861	GENERAL ELEMENTARY EDUCATION	NaN

	Text_1
344986	REGULAR INSTRUCTION
384803	CENTRAL
224382	NaN
305347	ADDL REGULAR PAY-NOT SMOOTHED
101861	REGULAR INSTRUCTION

[5 rows x 25 columns]

Which of these is a classification problem?

You'll continue your EDA in this exercise by computing summary statistics for the numeric data in the dataset. The data has been pre-loaded into a DataFrame called `df`.

You can use `df.info()` in the IPython Shell to determine which columns of the data are numeric, specifically type `float64`. You'll notice that there are two numeric columns, called `FTE` and `Total`.

- `FTE`: Stands for "full-time equivalent". If the budget item is associated to an employee, this number tells us the percentage of full-time that the employee works. A value of 1 means the associated employee works for the school full-time. A value close to 0 means the item is associated to a part-time or contracted employee.
- `Total`: Stands for the total cost of the expenditure. This number tells us how much the budget item cost.

After printing summary statistics for the numeric data, your job is to plot a histogram of the non-null `FTE` column to see the distribution of part-time and full-time employees in the dataset.

Instructions

- Print summary statistics of the numeric columns in the DataFrame `df` using the `.describe()` method.
- Import `matplotlib.pyplot` as `plt`.
- Create a histogram of the non-null `FTE` column. You can do this by passing `df['FTE'].dropna()` to `plt.hist()`.
- The title has been specified and axes have been labeled, so hit 'Submit Answer' to see how often school employees work full-time!

```
# Print the summary statistics
print(df.describe())

# Import matplotlib.pyplot as plt
import matplotlib.pyplot as plt

# Create the histogram
plt.hist(df['FTE'].dropna())

# Add title and labels
plt.title('Distribution of %full-time \n employee works')
plt.xlabel('% of full-time')
plt.ylabel('num employees')

# Display the histogram
plt.show()
```

Results :

<script.py> output:

	FTE	Total
count	449.000000	1.542000e+03
mean	0.493532	1.446867e+04
std	0.452844	7.916752e+04
min	-0.002369	-1.044084e+06
25%	NaN	NaN
50%	NaN	NaN
75%	NaN	NaN
max	1.047222	1.367500e+06

see : [img/graph1.svg](#)

Nice! The high variance in expenditures makes sense (some purchases are cheap some are expensive). Also, it looks like the FTE column is bimodal. That is, there are some part-time and some full-time employees.

Exploring datatypes in pandas

It's always good to know what datatypes you're working with, especially when the inefficient pandas type object may be involved. Towards that end, let's explore what we have.

The data has been loaded into the workspace as `df`. Your job is to look at the DataFrame attribute `.dtypes` in the IPython Shell, and call its `.value_counts()` method in order to answer the question below.

Make sure to call `df.dtypes.value_counts()`, and not `df.value_counts()`! Check out the difference in the Shell. `df.value_counts()` will return an error, because it is a Series method, not a DataFrame method.

How many columns with dtype object are in the data?

Possible Answers : => 2

- 2.
- 23.
- 64.
- 25.

```
df.dtypes.value_counts()
```

Results :

```
In [1]: df.dtypes.value_counts()
Out[1]:
object      23
float64      2
dtype: int64
```

Encode the labels as categorical variables

Remember, your ultimate goal is to predict the probability that a certain label is attached to a budget line item. You just saw that many columns in your data are the inefficient object type. Does this include the labels you're trying to predict? Let's find out!

There are 9 columns of labels in the dataset. Each of these columns is a category that has many possible values it can take). The 9 labels have been loaded into a list called LABELS. In the Shell, check out the type for these labels using `df[LABELS].dtypes`.

You will notice that every label is encoded as an object datatype. Because category datatypes are much more efficient your task is to convert the labels to category types using the `.astype()` method.

Note: `.astype()` only works on a pandas Series. Since you are working with a pandas DataFrame, you'll need to use the `.apply()` method and provide a lambda function called `categorize_label` that applies `.astype()` to each column, `x`.

Instructions

- Define the lambda function `categorize_label` to convert column `x` into `x.astype('category')`.
- Use the LABELS list provided to convert the subset of data `df[LABELS]` to categorical types using the `.apply()` method and `categorize_label`. Don't forget `axis=0`.
- Print the converted `.dtypes` attribute of `df[LABELS]`.

```
# Define the lambda function: categorize_label
categorize_label = lambda x: x.astype('category')

print(LABELS)
print('-----')
print(df[LABELS].dtypes)
print('-----')

# Convert df[LABELS] to a categorical type
df[LABELS] = df[LABELS].apply(categorize_label,axis=0)

# Print the converted dtypes
print(df[LABELS].dtypes)
```

Results :

```
<script.py> output:
['Function', 'Use', 'Sharing', 'Reporting', 'Student_Type', 'Position_Type', 'Object_Type', 'Pre_K', 'Operating_Status']
-----
Function      object
Use           object
Sharing       object
Reporting     object
Student_Type  object
Position_Type object
Object_Type   object
Pre_K         object
Operating_Status object
dtype: object
-----
Function      category
Use           category
Sharing       category
Reporting     category
Student_Type  category
Position_Type category
Object_Type   category
```


Pre_K	category
Operating_Status	category
dtype:	object

Great work! You're getting close to something you can work with. Keep it up!

Which of these is a classification problem?

As Peter mentioned in the video, there are over 100 unique labels. In this exercise, you will explore this fact by counting and plotting the number of unique values for each category of label.

The dataframe `df` and the LABELS list have been loaded into the workspace; the LABELS columns of `df` have been converted to category types.

`pandas`, which has been pre-imported as `pd`, provides a `pd.Series.nunique` method for counting the number of unique values in a Series.

Instructions

- Create the DataFrame `num_unique_labels` by using the `.apply()` method on `df[LABELS]` with `pd.Series.nunique` as the argument.
- Create a bar plot of `num_unique_labels` using `pandas`' `.plot(kind='bar')` method.
- The axes have been labeled for you, so hit 'Submit Answer' to see the number of unique values for each label.

```
# Import matplotlib.pyplot
import matplotlib.pyplot as plt

# Calculate number of unique values for each label: num_unique_labels
num_unique_labels = df[LABELS].apply(pd.Series.nunique)

# Plot number of unique values for each label
num_unique_labels.plot(kind='bar')

# Label the axes
plt.xlabel('Labels')
plt.ylabel('Number of unique values')

# Display the plot
plt.show()
```

Results :

Woah! That's a lot of labels to work with. How will you measure success with these many labels? You'll find out in the next video!

Penalizing highly confident wrong answers

As Peter explained in the video, log loss provides a steep penalty for predictions that are both wrong and confident, i.e., a high probability is assigned to the incorrect class.

Suppose you have the following 3 examples:

- A:y=1,p=0.85
- B:y=0,p=0.99
- C:y=0,p=0.51

Select the ordering of the examples which corresponds to the lowest to highest log loss scores. y is an indicator of whether the example was classified correctly. You shouldn't need to crunch any numbers!

Possible Answers : => 3

- Lowest: A, Middle: B, Highest: C.
- Lowest: C, Middle: A, Highest: B.
- Lowest: A, Middle: C, Highest: B.
- Lowest: B, Middle: A, Highest: C.

Results :

Yes! Of the two incorrect predictions, B will have a higher log loss because it is confident and wrong.

Computing log loss with NumPy

To see how the log loss metric handles the trade-off between accuracy and confidence, we will use some sample data generated with NumPy and compute the log loss using the provided function `compute_log_loss()`, which Peter showed you in the video.

5 one-dimensional numeric arrays simulating different types of predictions have been pre-loaded: `actual_labels`, `correct_confident`, `correct_not_confident`, `wrong_not_confident`, and `wrong_confident`.

Your job is to compute the log loss for each sample set provided using the `compute_log_loss(predicted_values, actual_values)`. It takes the predicted values as the first argument and the actual values as the second argument.

Instructions

- Using the `compute_log_loss()` function, compute the log loss for the following predicted values (in each case, the actual values are contained in `actual_labels`):
 - `correct_confident`.
 - `correct_not_confident`.
 - `wrong_not_confident`.
 - `wrong_confident`.
 - `actual_labels`.

```
# Compute and print log loss for 1st case
correct_confident = compute_log_loss(correct_confident, actual_labels)
print("Log loss, correct and confident: {}".format(correct_confident))

# Compute log loss for 2nd case
correct_not_confident = compute_log_loss(correct_not_confident, actual_labels)
print("Log loss, correct and not confident: {}".format(correct_not_confident))

# Compute and print log loss for 3rd case
wrong_not_confident = compute_log_loss(wrong_not_confident, actual_labels)
```

```
print("Log loss, wrong and not confident: {}".format(wrong_not_confident))

# Compute and print log loss for 4th case
wrong_confident = compute_log_loss(wrong_confident, actual_labels)
print("Log loss, wrong and confident: {}".format(wrong_confident))

# Compute and print log loss for actual labels
actual_labels = compute_log_loss(actual_labels, actual_labels)
print("Log loss, actual labels: {}".format(actual_labels))
```

Results :

```
In [1]: actual_labels
Out[1]: array([ 1.,  1.,  1.,  1.,  1.,  0.,  0.,  0.,  0.,  0.])

In [2]: correct_confident
Out[2]: array([ 0.95,  0.95,  0.95,  0.95,  0.95,  0.05,  0.05,  0.05,  0.05,  0.05])

<script.py> output:
Log loss, correct and confident: 0.05129329438755058
Log loss, correct and not confident: 0.4307829160924542
Log loss, wrong and not confident: 1.049822124498678
Log loss, wrong and confident: 2.9957322735539904
Log loss, actual labels: 9.99200722162646e-15
```

Wow! Log loss penalizes highly confident wrong answers much more than any other type. This will be a good metric to use on your models. You rock!

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :

Which of these is a classification problem?

Possible Answers : => 1

-

Results :
